



Introduction to Azure SQL

Module 1



Learning Units covered in this Module

- Lesson 1: Basic concepts of Azure SQL
- Lesson 2: Create an Azure SQL Database
- Lesson 3: Working with Managed Instances
- Lesson 4: Business Continuity Features

Lesson 1: Basic concepts of Azure SQL

Objectives

After completing this learning, you will be able to:

- Describe the basic concept and architecture
- Describe the difference between the purchase models.
- Describe the service tiers compute and hardware generation of the Azure SQL Database.



What is Microsoft Azure?

Microsoft Azure is Microsoft's public cloud computing platform

Over 150 countries across 60+ regions worldwide

Windows and Linux

Scale globally

- Reach more locations, faster, with the performance and reliability of a vast global infrastructure.

Safeguard data

- Rely on industry-leading data security in the region and across our network.

Promote sustainability


- Help build a clean-energy future and accelerate progress toward your sustainability goals.

[Azure Global Regions](#)

[Physical security of Azure datacenters](#)

[Microsoft Azure Sustainability](#)

What is Microsoft Azure?



[Azure global infrastructure experience \(microsoft.com\)](https://microsoft.com)

Cloud Hosting Models

Managed by customer

Managed by Microsoft

On-premises costs tend to be driven by hardware and data center management costs

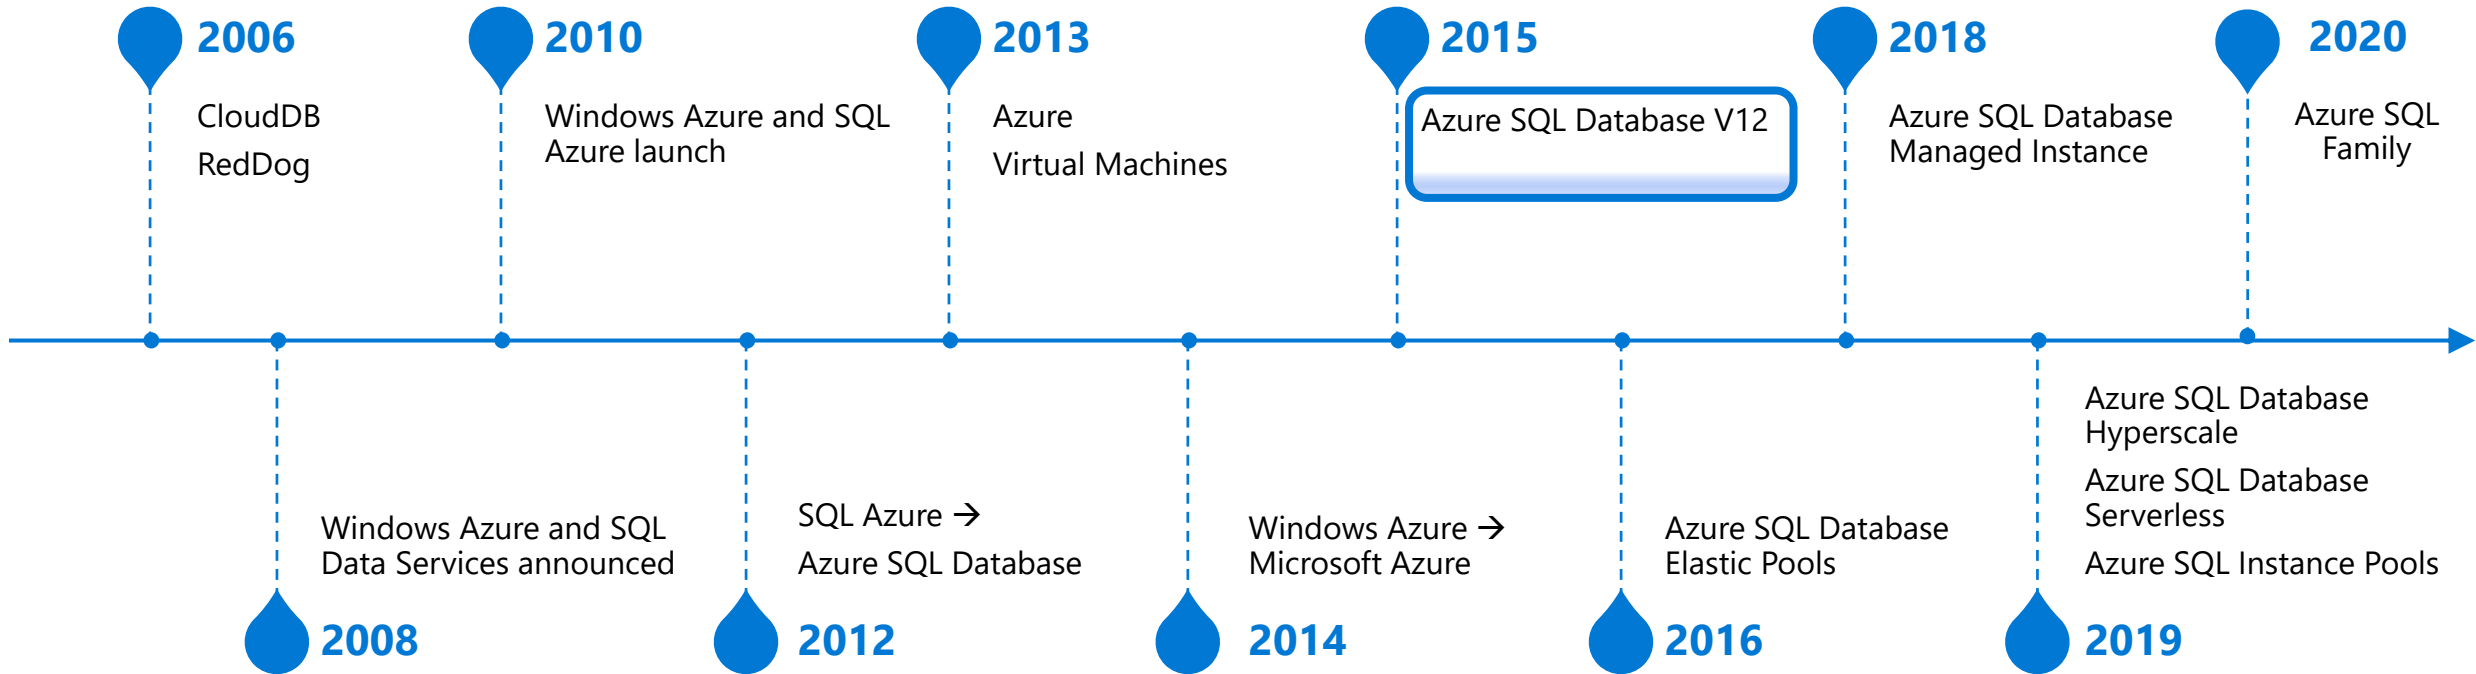
Infrastructure-as-a-Service reduces cost categories related to data center and compute

Platform-as-a-Service off-loads customers' most administrative tasks to Azure, further improving efficiency with machine-learning capabilities for performance and security

- **Managed Instance:** instance-level deployment for lift-shift existing apps to Azure, fully backward compatible
- **Single database:** database-level deployment for new apps

On-premises	Infrastructure (as a Service)	Platform (as a Service)
Applications	Applications	Applications
Data	Data	Data
High availability /DR/Backups	High availability /DR/Backups	High Availability/ DR/Backups
Database Provision/ Patch/Scaling	Database Provision/ Patch/Scaling	Database Provision/ Patch/Scaling
O/S provision /patching	O/S	O/S
Virtualization	Virtualization	Virtualization
Hardware	Hardware	Hardware
Datacenter Management	Datacenter Management	Datacenter Management
SQL Server 2017/2019	Azure SQL VMs	Azure SQL Database Azure SQL Managed Instance

Azure SQL has come a long way



Azure SQL Family

Explore Azure SQL database services



SQL Server on Azure Virtual Machines

Migrate your SQL workloads to Azure with ease while maintaining complete SQL Server compatibility and operating system-level access

[Learn more >](#)



Azure SQL Managed Instance

Modernize your existing SQL Server applications at scale with an intelligent, fully managed service

[Learn more >](#)






Azure SQL Database

Support modern cloud applications on an intelligent, fully managed service that includes serverless compute

[Learn more >](#)

Virtual Machine vs Managed Instance vs Azure SQL Database



A closer look... Best for

SQL Server in Azure VM

- Existing applications that require fast migration to the cloud with minimal changes or no changes.
- Teams that can configure, fine tune, customize, and manage high availability, disaster recovery, and patching for SQL Server.
- You need a customized environment with full administrative rights.
- SQL Server instances with up to 64 TB of storage. The instance can support as many databases as needed.

Azure SQL Managed Instance

- New applications or existing on-premises applications that want to use the latest stable SQL Server features and that are migrated to the cloud with minimal changes.
- Teams that need built-in high availability, disaster recovery, and upgrade for the database.
- Teams that do not want to manage the underlying operating system and configuration settings.
- Databases of up to 8 TB, or larger databases that can be horizontally or vertically partitioned using a scale-out pattern.

Azure SQL Database

- New cloud-designed applications that want to use the latest stable SQL Server features and have time constraints in development and marketing.
- Teams that need built-in high availability, disaster recovery, and upgrade for the database.
- Teams that do not want to manage the underlying operating system and configuration settings.
- Databases of up to 4 TB, or larger databases that can be horizontally or vertically partitioned using a scale-out pattern.

Azure SQL PaaS deployment options



Azure SQL Database

Managed instances

Best for most lift-and-shift migrations to the cloud



Single instance

- SQL Server surface area (vast majority).
- Native virtual network support.
- Fully managed service.

Instance pool*

- Resource sharing between multiple instances to price optimize.
- Simplified performance management for multiple databases.
- Fully managed service.

*Instance pool is currently in public preview, more details [here](#)

Databases

Best for modern cloud applications. Hyperscale and serverless options are available



Single database

- Hyperscale storage (up to 100TB).
- Serverless compute.
- Fully managed service.


Elastic pool

- Resource sharing between multiple databases to price optimize.
- Simplified performance management for multiple databases.
- Fully managed service.

Azure SQL Database Purchasing models and resources


Database Transaction Unit (DTU) model

- Pre-packaged, bundled unit that represents the database power.
- Designed for **predictable performance**, but somewhat inflexible and limited in options.
- DTU sizing offers simplicity of choice.
- Blended measure of CPU, memory, and read-write rates.



vCore model

- À La carte approach deconstructs the DTU model into separate parts.
- Customers can select compute and storage independently.
- Allows customers to right-size their compute requirements in the cloud.
- vCore sizing offers flexibility of choice.
- Provides a choice between a **provisioned compute tier** and a **serverless compute tier**.

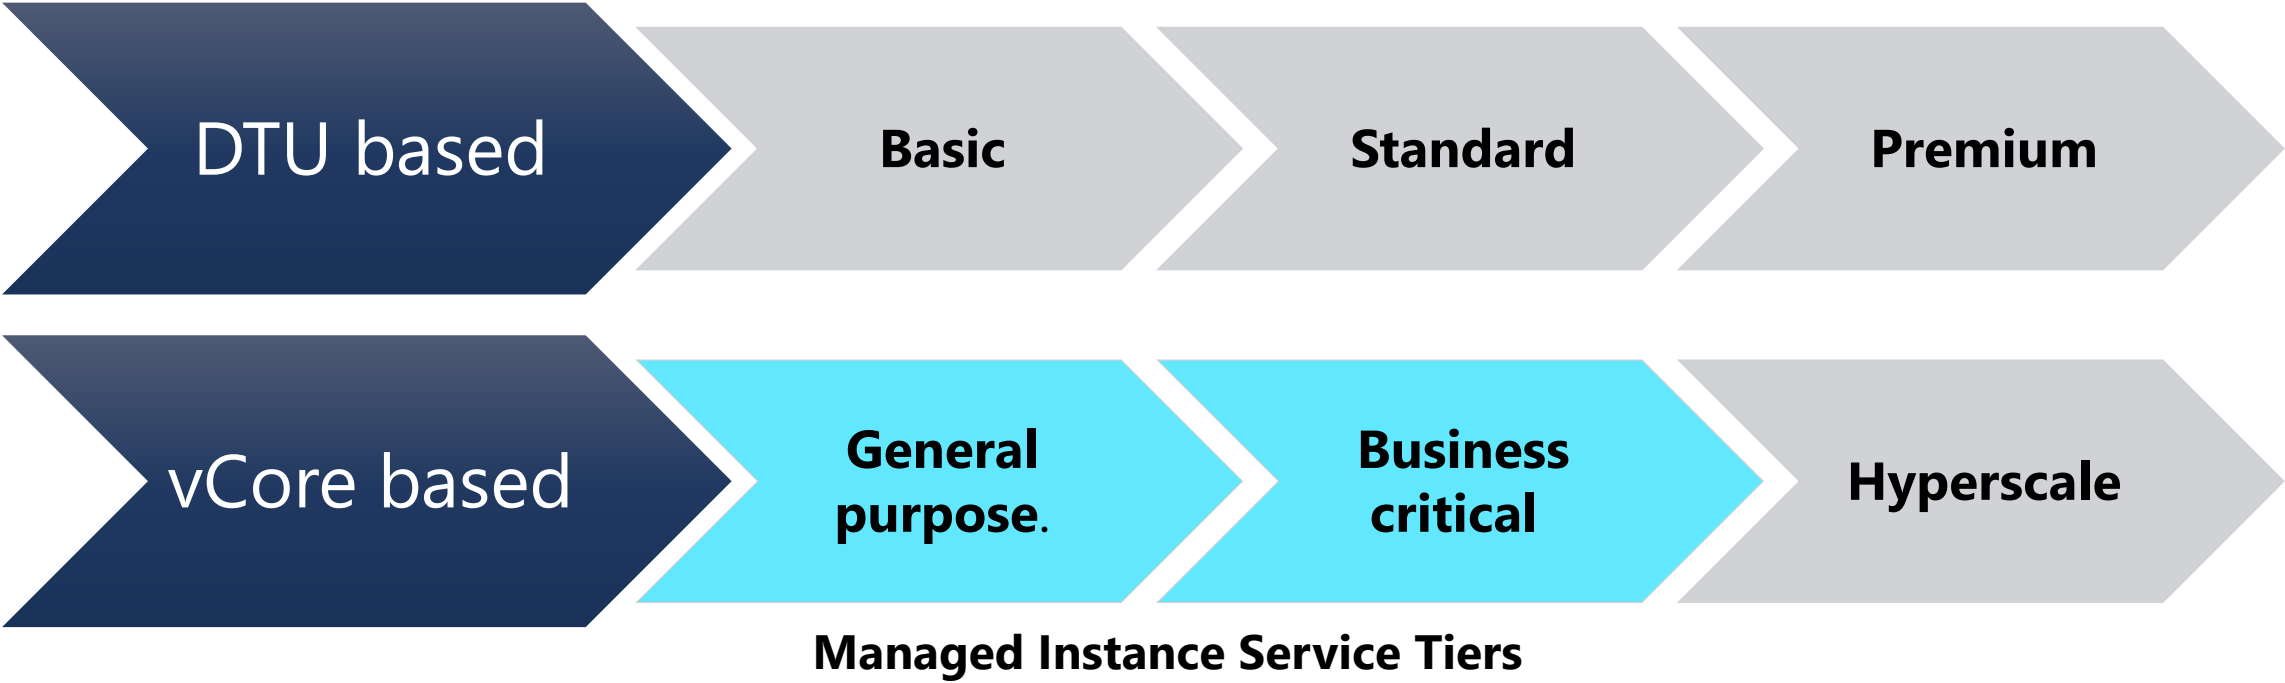


Purchasing models available for Azure SQL Database deployment:

Purchasing model/ Deployment model	DTU based	vCore based
Single Database	✓	✓
Elastic Pool	✓	✓
Managed Instance	✗	✓
Hyperscale service tier (single databases)	✗	✓

Overview Service Tiers

Standard 100DTUs = 1vCore
Premium 125DTUs = 1vCore



Changing Performance Levels (DTU)

PowerShell

- Set-AzSqlDatabase

REST

- Update database

Azure CLI

- az sql db update

T-SQL

- ALTER DATABASE ... MODIFY
(EDITION = ...)

Service and compute tier

Select from the available tiers based on the needs of your workload. The vCore model provides a wide range of configuration controls and offers Hyperscale and Serverless to automatically scale your database based on your workload needs. Alternately, the DTU model provides set price/performance packages to choose from for easy configuration. [Learn more](#)

Service tier

DTUs [Compare DTU options](#)

5 (Basic)

Data max size (GB)

DTU-based purchasing model

- Basic (For less demanding workloads)
- Standard (For workloads with typical performance requirements)
- Premium (For IO-intensive workloads)

Changing Performance Levels (vCore)

PowerShell

- Set-AzSqlDatabase

REST

- Update database

Azure CLI

- az sql db update

T-SQL

- ALTER DATABASE ... MODIFY (EDITION = ...)

Service and compute tier

Select from the available tiers based on the needs of your workload. The vCore model provides a wide range of configuration controls and offers Hyperscale and Serverless to automatically scale your database based on your workload needs. Alternately, the DTU model provides set price/performance packages to choose from for easy configuration. [Learn more](#)

Service tier

General Purpose (Scalable compute and storage options) ▼

Compute tier

vCore-based purchasing model

General Purpose (Scalable compute and storage options)

Hyperscale (On-demand scalable storage)

Business Critical (High transaction rate and high resiliency)

DTU-based purchasing model

Basic (For less demanding workloads)

Standard (For workloads with typical performance requirements)

Premium (For IO-intensive workloads)

up to 80 vCores, up to 408 GB memory

[Change configuration](#)

Compute Hardware

Select the hardware configuration based on confidential computing hardware depends on

Hardware Configuration

Save money

Already have a SQL Server License? Save with a license you already own with Azure Hybrid Benefit. Actual savings may vary based on region and performance tier. [Learn more](#)

☐ Yes ☒ No

vCores [Compare vCore options](#)



2

Data max size (GB) ⓘ



32

DTU-based model: Service Tiers

	Basic	Standard									Premium					
		S0	S1	S2	S3	S4	S6	S7	S9	S12	P1	P2	P4	P6	P11	P15
Built for	Light transactional workloads	Medium transactional workloads									Heavy transactional workloads					
Available SLA	99.99%															
Database max. size	2 GB	250 GB			1 TB										4 TB	
Point-in-time restore ("oops" recovery)	Any point within 7 days	7-35 days (7 days by default)														
Business continuity	Active geo-replication, up to four readable secondary backups															
Security	Auditing, row-level security, dynamic data masking, Advanced Threat Protection															
Performance objectives	Transactions per hour	Transactions per minute									Transactions per second					
Database transaction units (DTUs)	5	10	20	50	100	200	400	800	1,600	3,000	125	250	500	1,000	1,750	4,000

\$

\$\$\$

[Resource limits for single databases using the DTU-based purchasing model](#)

Changing Compute Tier and Hardware (vCore)

Service and compute tier

Select from the available tiers based on the needs of your workload. The vCore model provides a wide range of configuration controls and offers Hyperscale and Serverless to automatically scale your database based on your workload needs. Alternately, the DTU model provides set price/performance packages to choose from for easy configuration. [Learn more](#)

Service tier

General Purpose (Scalable compute and storage options) ▼

[Compare service tiers](#) ↗

Compute tier



Provisioned - Compute resources are pre-allocated. Billed per hour based on vCores configured.



Serverless - Compute resources are auto-scaled. Billed per second based on vCores used.

Compute Hardware

Select the hardware configuration based on your workload requirements. Availability of compute optimized, memory optimized, and confidential computing hardware depends on the region, service tier, and compute tier.




Hardware Configuration

Gen5

up to 80 vCores, up to 408 GB memory

[Change configuration](#)

vCore-based purchasing model: Provisioned compute tier

Service tier	 General purpose		 Business critical		 Hyperscale
Best for	Most budget-oriented workloads		Critical business applications with high IO requirements.		VLDB OLTP and HTAP workloads with highly scalable storage and read-scale requirements.
Deployment option	Single / Elastic Pools Managed Instance		Single / Elastic Pools Managed Instance		Single
Compute tiers	Gen5: 2 to 80 vCore Fsv2-series*: 72 vCore Serverless: 0.5 to 16 vCore		Gen5: 4 to 80 vCore	Gen5: 2 to 80 vCore M-series*: 128 vCore	Gen5: 4 to 80 vCore
Storage	Premium remote		Local SSD		Local SSD Cache
	5GB – 4TB per instance	32GB – 16TB per instance	5GB – 4TB per instance	32GB – 16TB per instance	Scale from 5GB to 100TB of storage in 1GB increments
In-Memory	Not supported		Supported		Not supported
Read-write IO	~2ms for all data access		<0.5ms for all data access		<0.5ms for hot data access ~2ms otherwise
Availability	1 replica, no read-scale replicas		3 read replicas, zone-redundant HA	3 replicas, 1 read-scale replica, zone-redundant HA	Primary read/write replica + up to 4 read replicas
Backups	RA-GRS, 7-35 days (7 days by default)		RA-GRS, 7-35 days (7 days by default)		LRS, ZRS, RA-GRS, 7-35 days (7 days by default)

For latest information reference: <https://azure.microsoft.com/en-us/pricing/details/sql-database/>

* Fsv2-series and M-series are currently in preview.

Choose from hardware generations

	Gen 5	Fsv2-series	M-series	DC-Series
Hardware	Intel E5-2673 v4 (Broadwell) 2.3 GHz processors, fast eNVM SSD vCore=1 LP (hyper-thread)	Intel Xeon Platinum 8168 (SkyLake) processors	Intel Xeon E7-8890 v3 2.5 GHz processors	Intel® XEON E-2288G processors. Featuring Intel Software Guard Extension (Intel SGX))
Performance levels	1 to 80 vCores	72 vCores (1 vCore = 1 hyper-thread)	128 vCores (1 vCore = 1 hyper-thread)	Provision up to 8 vCores (physical)
Memory	5.1 GB per vCore for provisioned compute Up to 24 GB per vCore for serverless compute	1.9 GB per vCore	29 GB per vCore	4.5 GB per vCore
Storage	Up to 4TB remote SSD storage Up to 4TB local SSD storage	Up to 4TB remote SSD storage	Up to 4TB local SSD storage	Up to 4TB remote SSD storage

- Balance performance requirements and price with two hardware generations
- Match your on-premises application behavior
- Fsv2, M, and DC series are Generally available but in limited regions

Provisioned compute and serverless meet different needs

Optimize compute provisioning and billing for your workload

Databases with provisioned compute...

Provision compute resources upfront.

Bill on an hourly basis.

Common scenarios

Workloads with regular and substantial compute utilization.

Multiple databases with bursty usage patterns that can be consolidated into a single server and use *elastic pools* for better price optimization.



Serverless databases...

Scale up or down to meet workload requirements, instead of pre-provisioning.

Bill on a per-second basis.

Common scenarios

Workloads with unpredictable and intermittent usage patterns or performance requirements .

Workloads where the requirements are unknown, and you can delegate compute sizing to the service.



Serverless price to performance with per-second billing

Compute resources scale dynamically up or down based on workload requirements.


Configure minimum and maximum vCores to define the range of available compute capacity.

Use auto-pause delay to define the time period the dataset must be inactive before pausing.

The database is automatically resumed when the next login or other activity occurs.

Pay for compute based on the vCores and memory used per second, with lowest billing based on configured vCore minimum.

When the database is paused, the compute cost is zero and only storage costs are incurred.



Lesson 2: How to create Azure SQL Database


Objectives

After completing this learning, you will be able to:

- Know which prerequisites are needed before you can create an Azure SQL Database.
- Create your first Azure SQL Database.



Azure SQL Database Server is not a machine...



*TDS = Tabular Data Stream

Server Provisioning

Service head that contains databases

Connect via Fully Qualified Domain Name

Initially contains only a **master** database

Create SQL Database Server

Microsoft

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription *

PFE Subscription

Resource group *

(New) AzureSQLDatabaseRG

[Create new](#)

Server details

Enter required settings for this server, including providing a name and location.

Server name *

dbssql essentials

.database.windows.net

Location *

(US) East US

Authentication

Select your preferred authentication methods for accessing this server. Create a server admin login and password to access your server with SQL authentication, select only Azure AD authentication [Learn more](#) using an existing Azure AD user, group, or application as Azure AD admin [Learn more](#), or select both SQL and Azure AD authentication.

Authentication method

☒

Use SQL authentication

☐

Use only Azure Active Directory (Azure AD) authentication

☐

Use both SQL and Azure AD authentication

Server admin login *

sqlserveradmin

Password *

.....

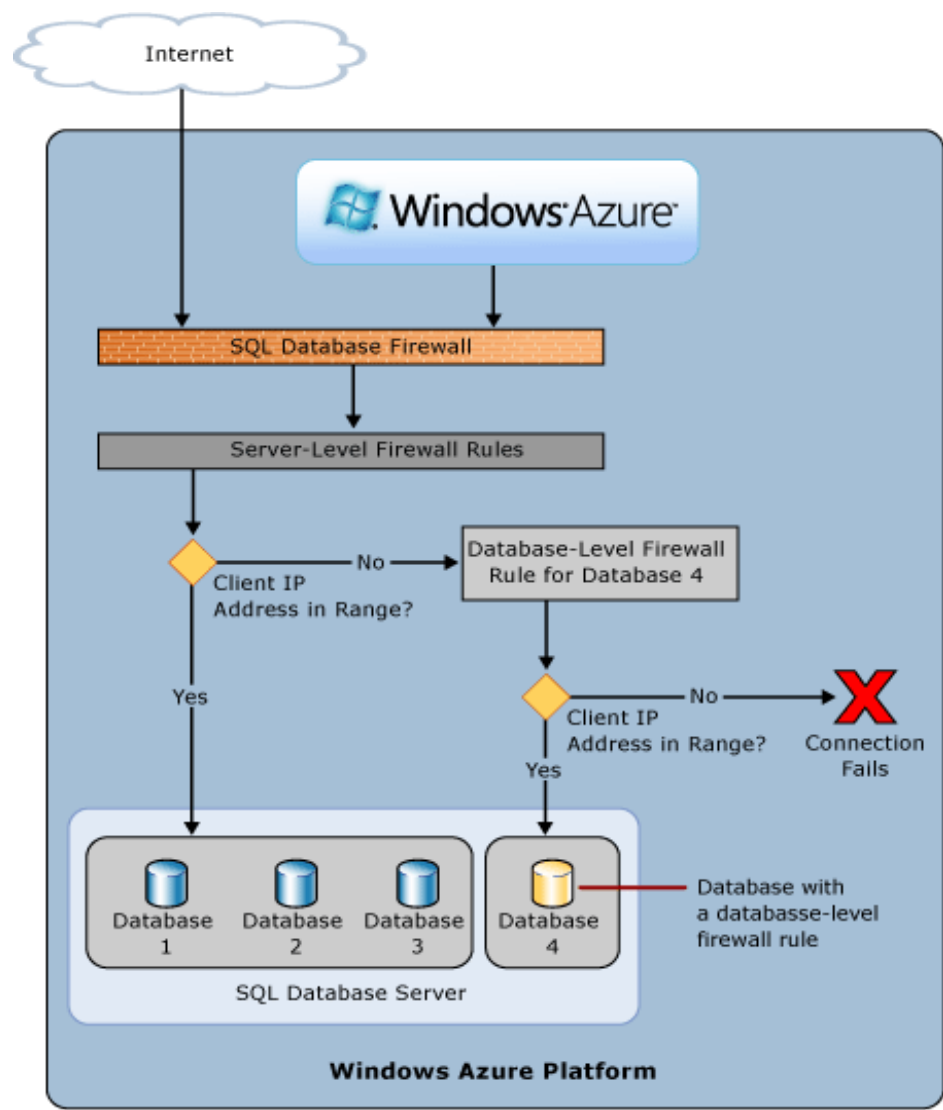
Confirm password *

.....

[Review + create](#)

[Next : Networking >](#)

Azure SQL Database Firewall



IP Address-based access control for SQL Database

Rules at the server and/or database level

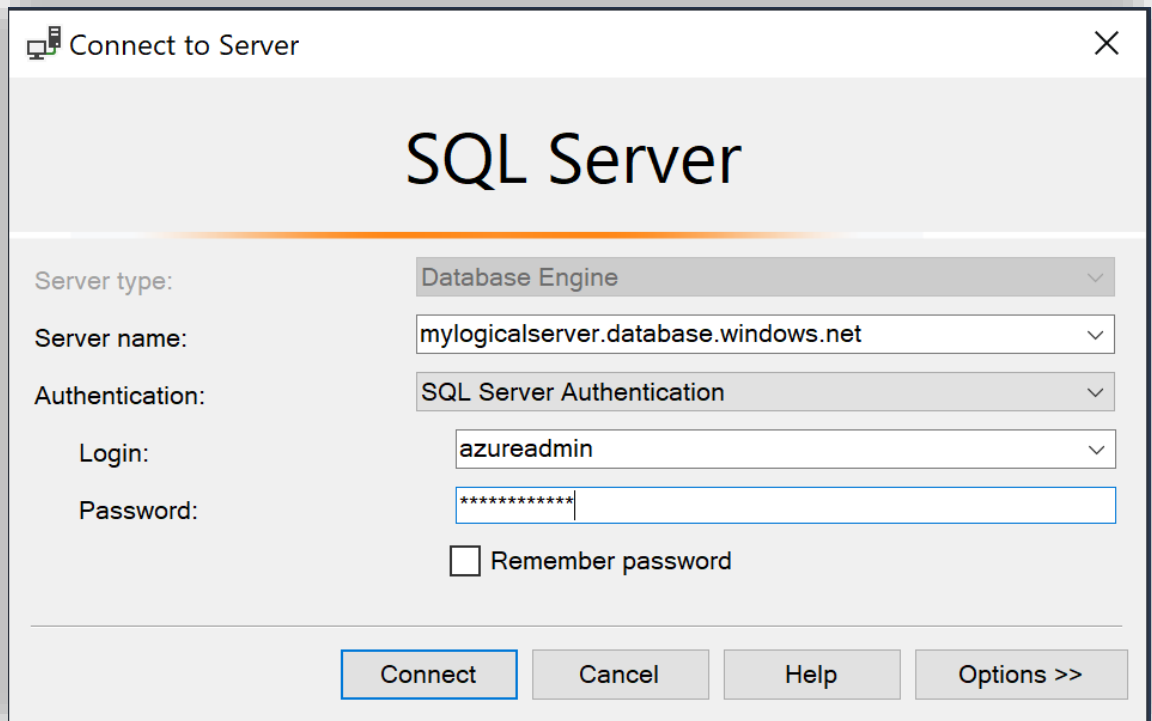
No IP authorized by default, not even Azure itself

SQL Server Management Studio

Download
the latest
version of
SSMS.

Get the fully
qualified
domain
name of
your Azure
SQL Server.


Connect to
your SQL
Database.



The screenshot shows the 'Connect to Server' dialog box with the following fields and options:

- Server type:** Database Engine
- Server name:** mylogicalserver.database.windows.net
- Authentication:** SQL Server Authentication
- Login:** azureadmin
- Password:** (masked with asterisks)
- ☐ Remember password
- Buttons:** Connect, Cancel, Help, Options >>

SQL Server Data Tools



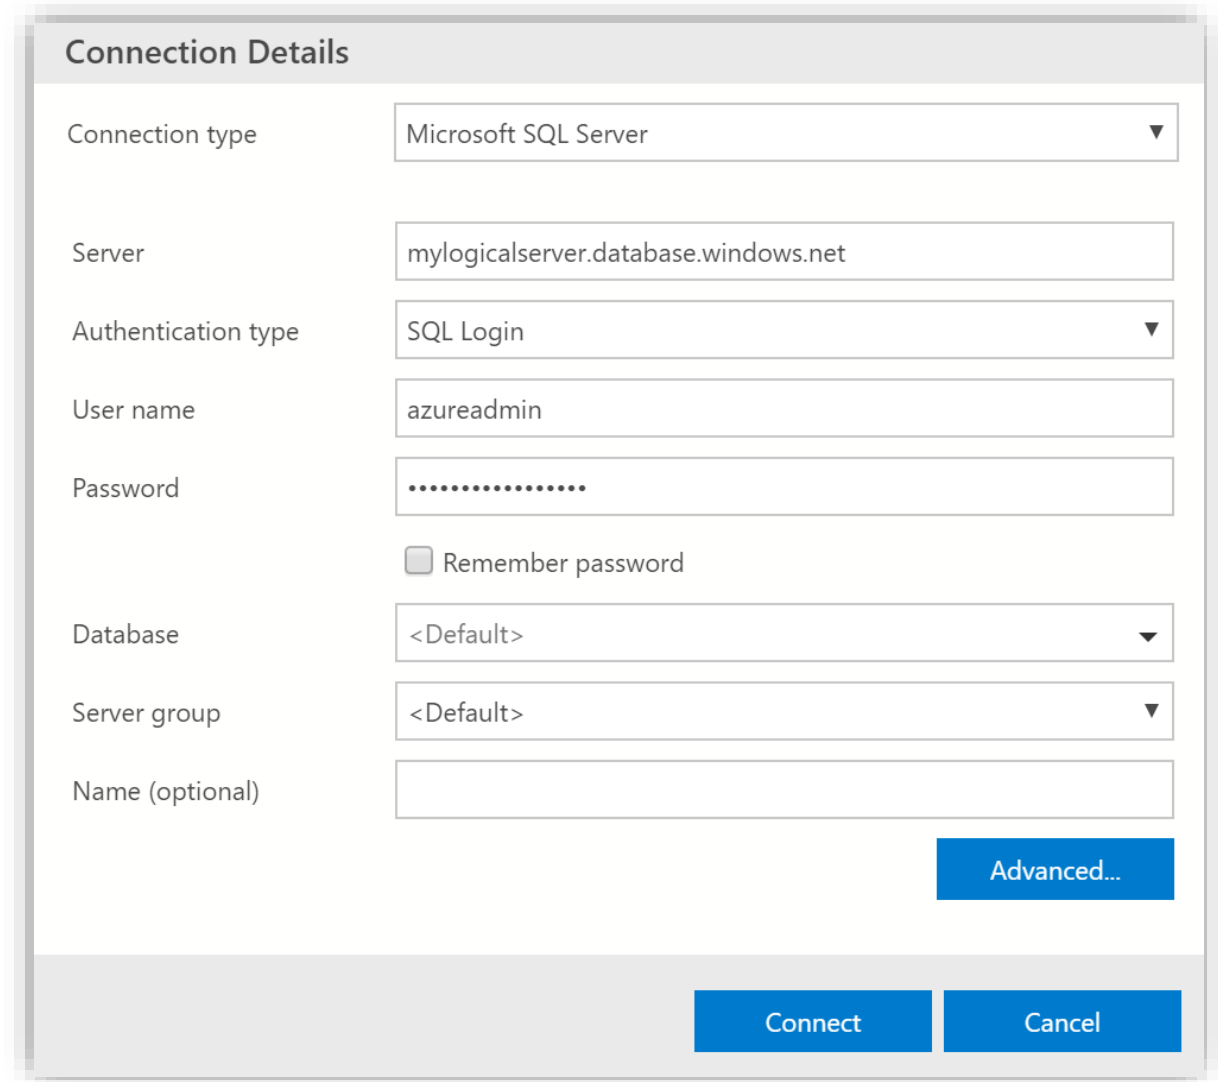
The screenshot shows the 'Connect' dialog box in SQL Server Data Tools. The 'Browse' tab is selected, showing a search bar and a tree view with 'Local', 'Network', and 'Azure' options. Below the tree view, the 'Server Name' field is populated with 'mylogicalserver.database.windows.net'. The 'Authentication' dropdown is set to 'SQL Server Authentication'. The 'User Name' field contains 'azureadmin', and the 'Password' field is masked with dots. There is an unchecked 'Remember Password' checkbox. The 'Database Name' dropdown is set to '<default>'. At the bottom right, there is an 'Advanced...' link and 'Connect' and 'Cancel' buttons.

Azure Data Studio

Download the latest version of Azure Data Studio.

Get the fully qualified domain name of your Azure SQL Server.

Connect to your SQL Database.




The screenshot shows the 'Connection Details' dialog box in Azure Data Studio. It contains the following fields and options:

- Connection type:** A dropdown menu set to 'Microsoft SQL Server'.
- Server:** A text box containing 'mylogicalserver.database.windows.net'.
- Authentication type:** A dropdown menu set to 'SQL Login'.
- User name:** A text box containing 'azureadmin'.
- Password:** A text box with masked characters (dots). Below it is a checkbox labeled 'Remember password' which is currently unchecked.
- Database:** A dropdown menu set to '<Default>'.
- Server group:** A dropdown menu set to '<Default>'.
- Name (optional):** An empty text box.

At the bottom right of the dialog, there are three buttons: 'Advanced...' (in a blue box), 'Connect' (in a blue box), and 'Cancel' (in a blue box).

Azure Portal



Demonstration

Create your first Azure SQL Database

- Create a new server, configure the firewall rules and create an Azure SQL Database (DTU based).



Lesson 3: Working with Managed Instances

Azure SQL Managed Instance



SQL Server on Azure Virtual Machines

Best for lift and shift and workloads requiring OS access

Infrastructure-as-a-Service



Azure SQL Managed Instance

Best for modernizing existing apps

Platform-as-a-Service



Azure SQL Database

Best for supporting modern cloud apps




Azure SQL Edge

Best for extending apps to IoT edge


Edge Computing

Which Azure SQL offering is right for you?



Azure SQL Managed Instance Service Tiers

Service tier


Select from the latest vCore service tiers available for Azure SQL Managed Instance including General Purpose and Business Critical. [Learn more](#) 

Service tier ⓘ

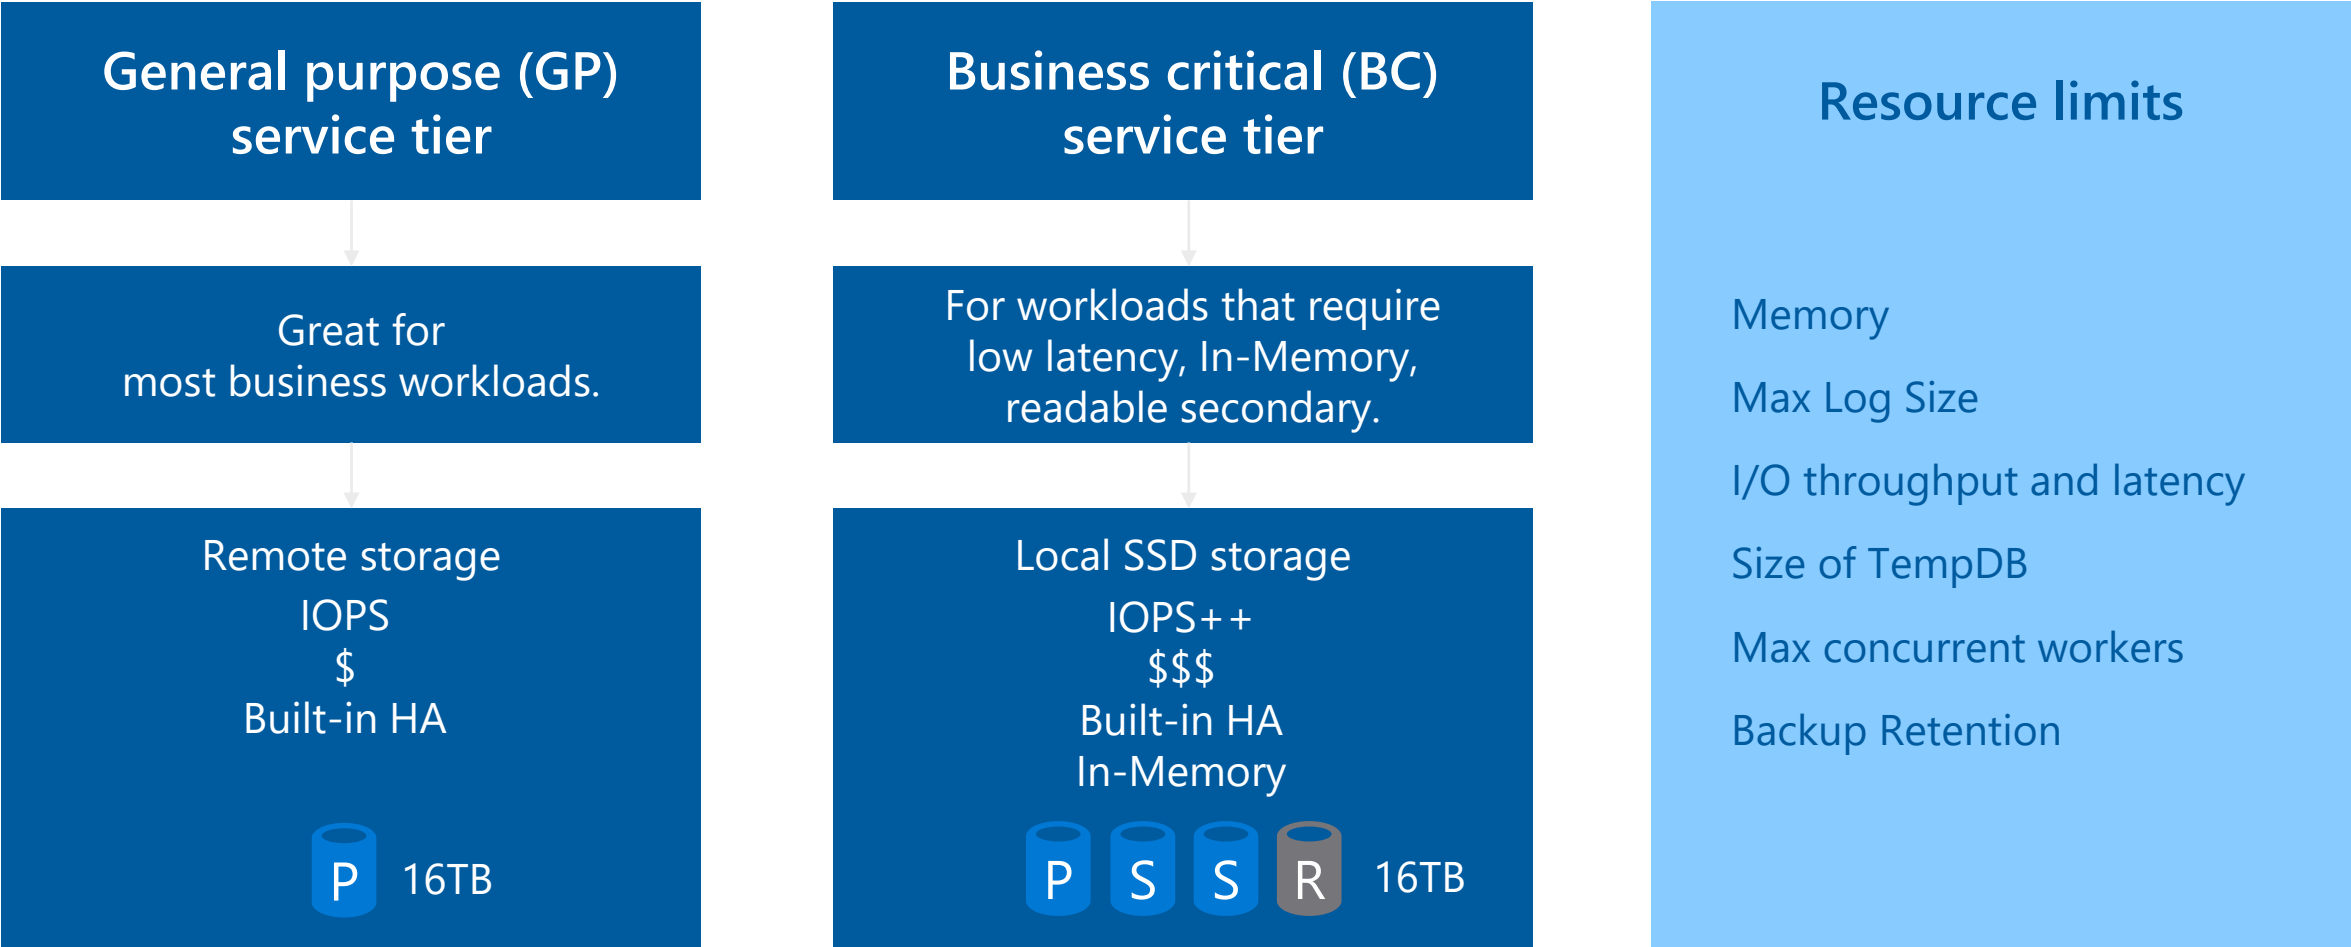
- ☒ General Purpose (4-80 vCores, 32 GB-16 TB storage capacity, Fast storage) - for most production workloads
- ☐ Business Critical (4-80 vCores, 32 GB-4 TB storage capacity, Super fast storage) - for IO-intensive and compute-intensive workloads

Next-gen General Purpose (preview)


Enabled **Disabled**

i Next-gen General Purpose tier is currently in preview. By using this set of preview features, you confirm that you agree that your use of these features is subject to the preview terms in the agreement under which you obtained Microsoft Azure Services. [Learn more](#) 

Azure SQL Managed Instance Service Tiers




Azure SQL Managed Instance Service Tiers



Next-gen General Purpose (Preview)

Announced March 2024


Select from the latest vCore service tiers available for Azure SQL Managed Instance including General Purpose and Business Critical. [Learn more](#) 

Service tier ⓘ

- ☒ General Purpose (4-80 vCores, 32 GB-16 TB storage capacity, Fast storage) - for most production workloads
- ☐ Business Critical (4-80 vCores, 32 GB-4 TB storage capacity, Super fast storage) - for IO-intensive and compute-intensive workloads

Next-gen General Purpose (preview)

Enabled **Disabled**

i Next-gen General Purpose tier is currently in preview. By using this set of preview features, you confirm that you agree that your use of these features is subject to the preview terms in the agreement under which you obtained Microsoft Azure Services. [Learn more](#) 

Next-gen General Purpose (Preview)

Announced March 2024

Designed for businesses with higher performance requirements while offering the same baseline cost as the General Purpose service tier

Significant upgrades to performance, scalability, and resource flexibility over the General Purpose service tier

Uses managed disks instead of page blobs, which drastically improve storage performance metrics

3 free IOPS for every GB of reserved storage

Support of up to 500 databases per instance, and a max storage size of 32 TB

SQL Managed Instance Hardware Generations

Standard-series (Gen 5)	Premium-series	Premium-series Mem Optimized
<ul style="list-style-type: none">• CPU: Intel Broadwell, Skylake and Cascade Lake, 2.3-2.5 GHz• vCore range: 4 – 80• Memory / vCore: 5.1 GB• Max instance memory: 408 GB• Max instance storage<ul style="list-style-type: none">• General Purpose: 16 TB• Business Critical: 4 TB	<ul style="list-style-type: none">• CPU: Latest 3rd Gen Intel 8370C (Ice Lake), 2.8 GHz• vCore range: 4 – 80• Memory / vCore: 7 GB• Max instance memory: 560 GB• Max instance storage<ul style="list-style-type: none">• General Purpose: 16 TB• Business Critical: 5.5 TB	<ul style="list-style-type: none">• CPU: Latest 3rd Gen Intel 8370C (Ice Lake), 2.8 GHz• vCore range: 4 - 64• Memory / vCore: 13.6 GB• Max instance memory: 870 GB• Max instance storage<ul style="list-style-type: none">• General Purpose: 16 TB• Business Critical: 16 TB

Deploy, Connect, Configure



Deploy

Choose region, service tier, vCores, and Max Storage
You choose Max Storage for instance based on possible max size
Choose hardware options NEW
Service-aided subnet configuration



Connect

Inside the virtual network with Private IP
Outside using public endpoint/port with NSG



Configure

Configure outside of SQL with the portal or CLI
Change tier, vCores, max storage, maintenance windows NEW
Configure inside SQL with T-SQL or SSMS



What's different?

No OS or file system access
No manual restarts or SQL Server Config Manager
Tempdb managed
Some global trace flags allowed



Evergreen version of SQL



Get rid of perpetual SQL
major version upgrades

This is the LAST SQL
Server upgrade that
you ever perform



Newest / Azure-only
features

Leverage latest SQL
Server innovation
available as “cloud-first”


NEW: Database format alignment with SQL Server 2022

Lesson 4: Business Continuity Features

Objectives

After completing this learning, you will be able to:

- Understand the various business continuity options within Azure SQL Database.
- Understand how to copy and export Azure SQL Databases.
- Understand how to perform a point-in-time restore.
- Understand how to perform a restore of a deleted database.



Business Continuity Problem


Mechanisms, Policies, and Procedures that enable a business to continue operating in face of disruption.

Customer Error



Customer deleted their data by mistake.

Equipment Failures



Hardware or other equipment failures

Natural Disasters



Flood, Hurricanes, Tornadoes, Earthquakes

Manmade Disasters



Wars or major ecological accidents


Availability vs High Availability vs Disaster Recovery



Availability

Every database comes with core resiliency and availability, that protects it against software or hardware failures.

Industry leading financially backed SLA of 99.99% availability.



High Availability

Continuous availability of the database provided through Zone Redundancy.

Automatic in region recovery from zonal hardware and software failures that's transparent to applications.

Higher SLA of 99.995% availability.



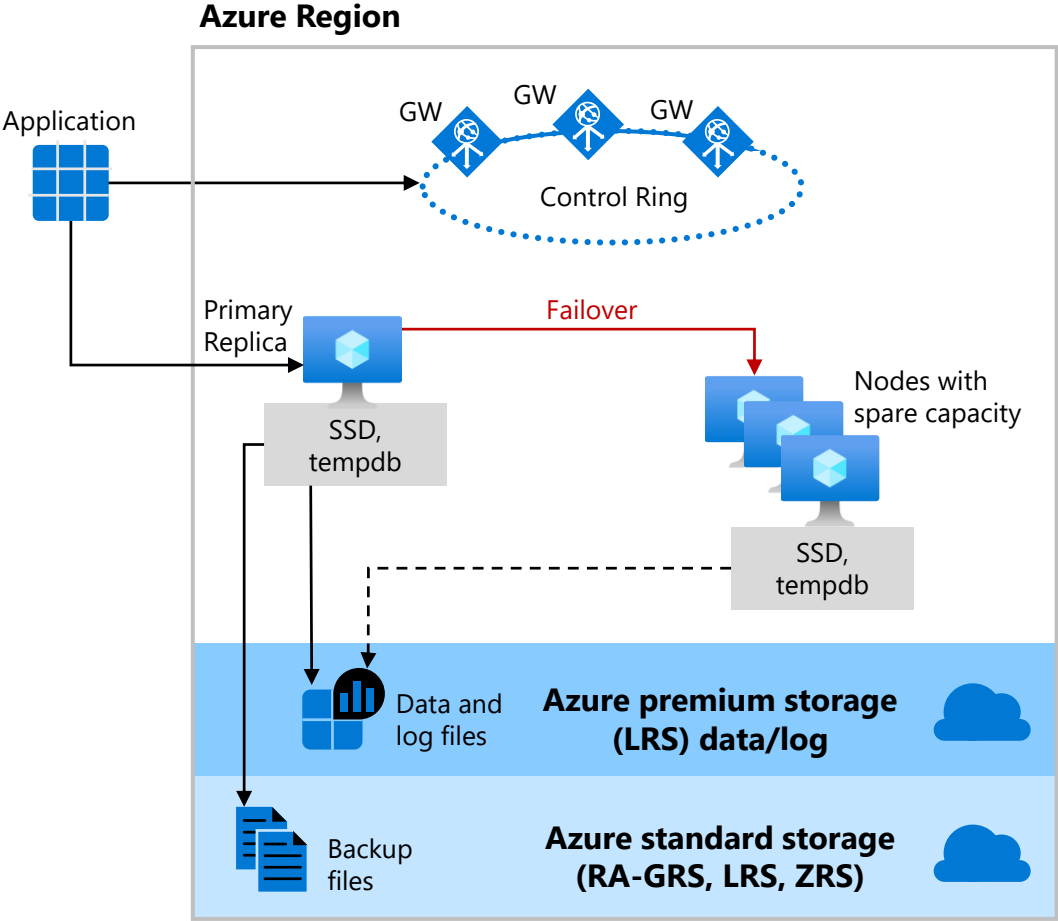
Disaster Recovery

Customer configured database replicas or failover groups

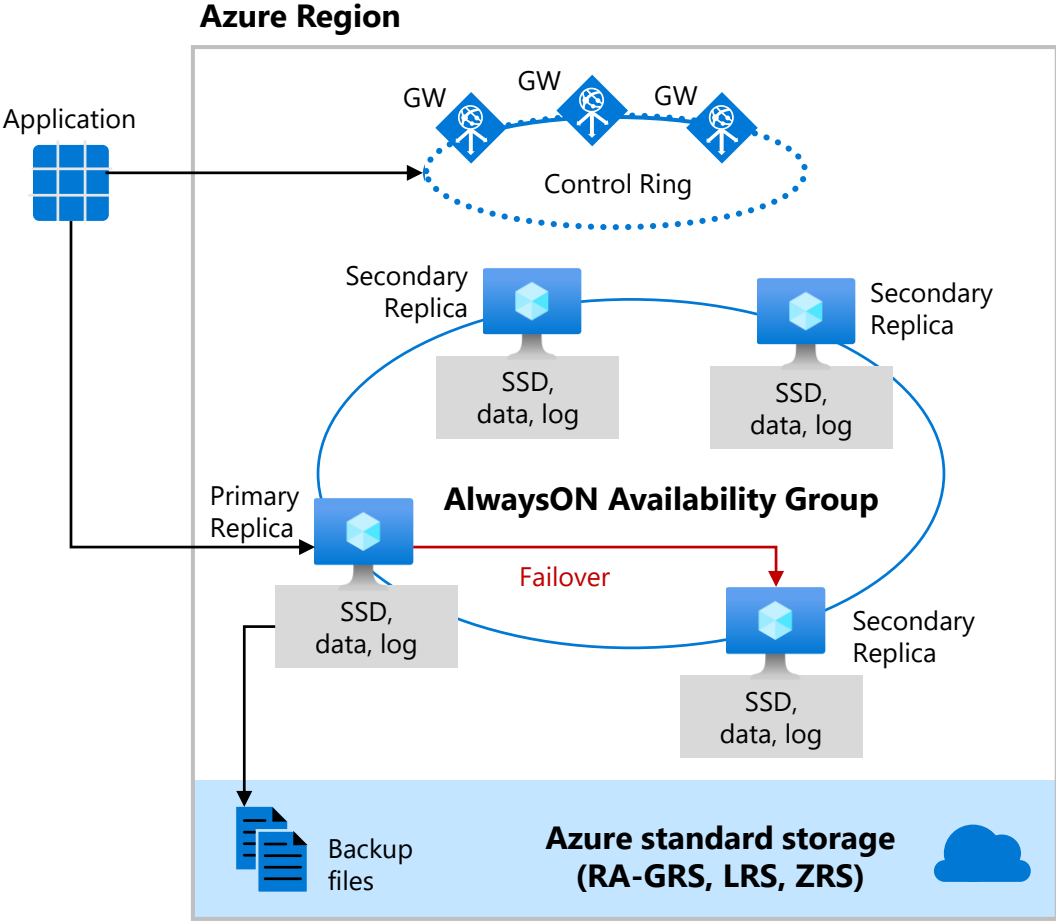
Ability to quickly recover the database from a catastrophic regional failure to provide business continuity.

Protect your business by maximizing application availability.

High Availability: General Purpose vs Business Critical



Basic (DTU), Standard (DTU), General Purpose (vCore)



Premium (DTU) and Business Critical (vCore)

High Availability - General Purpose

Basic (DTU), Standard (DTU), General Purpose (vCore)

Behaves like Failover Cluster Instance

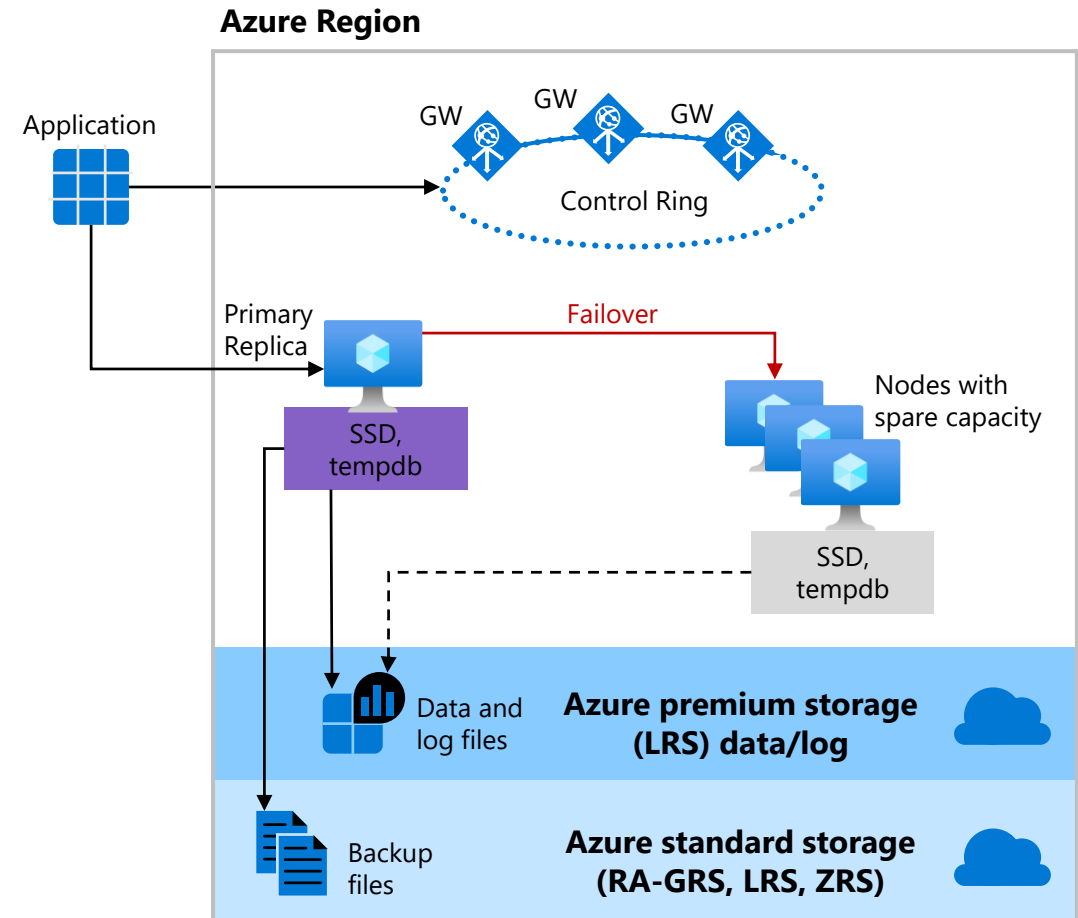
Remote storage provides data redundancy

Backup files are in a different location with geo-redundancy

Failover decisions based on SQL and Service Fabric

Recovery time depends on spare capacity

Connectivity redirection built-in



High Availability - Business Critical

Premium (DTU) and Business Critical (vCore)

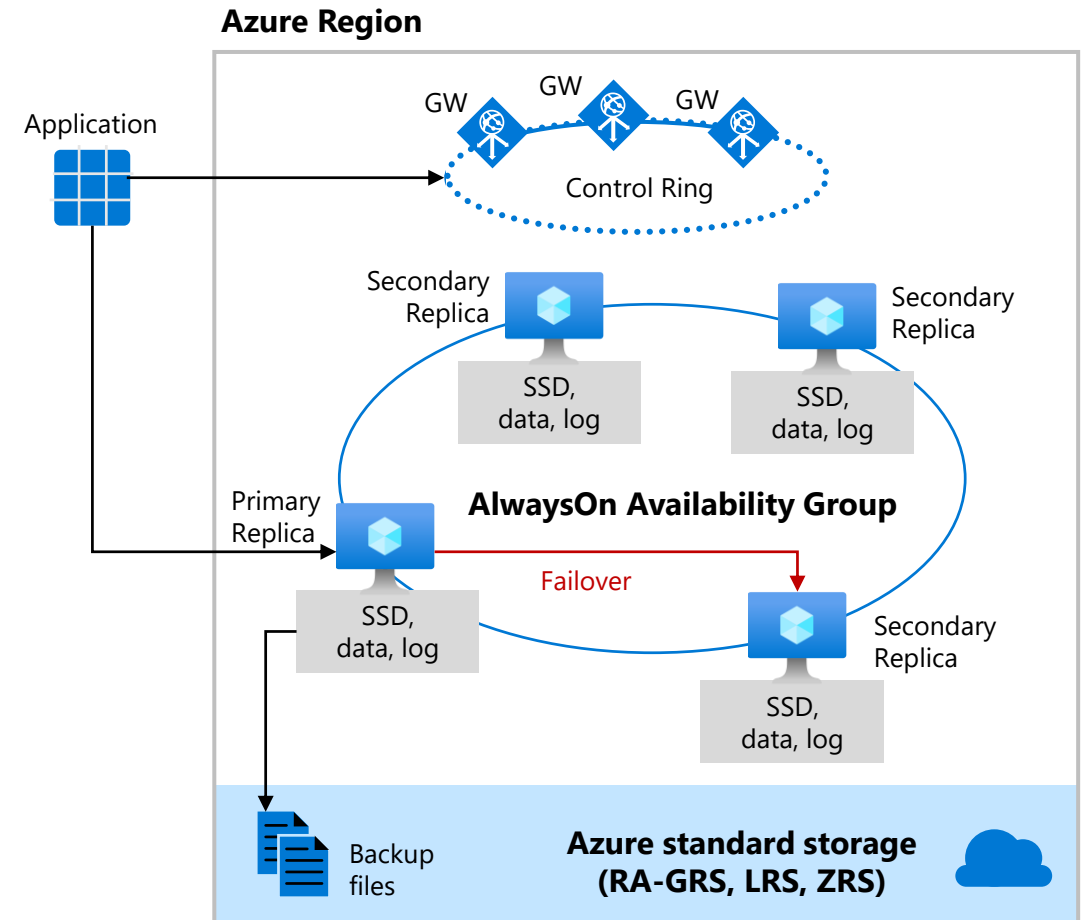
High availability is implemented using a technology like SQL Server Always On Availability Groups.

High availability is achieved by replicating both compute and storage to additional nodes.

The cluster includes a single primary replica for read-write workloads, and up to three secondary replicas.


The failover is initiated by the Azure Service Fabric.

As an extra benefit, the premium availability model includes Read Scale-Out feature.



Backup storage redundancy

To enable high durability of backups several ways of replication are offered on instance creation.



The backups can be all located within

1. LRS: The same building (Local)
2. ZRS: Same region, different buildings (Zone)
3. GRS: Across paired regions (Geo)
4. GZRS: Different buildings AND paired regions (Geo-Zone)

Zone redundant configuration – Premium and Business Critical


By default, the cluster of nodes for the premium availability model is created in the same datacenter.

SQL Database can place different replicas of the Business-Critical database to different availability zones in the same region.

The routing is controlled by Azure Traffic Manager (ATM).

The zone redundant databases have replicas in different datacenters with some distance between them, the increased network latency may impact the performance.

Zone redundant configurations are available in the Premium, Business Critical, General Purpose and Hyperscale Service Tiers.



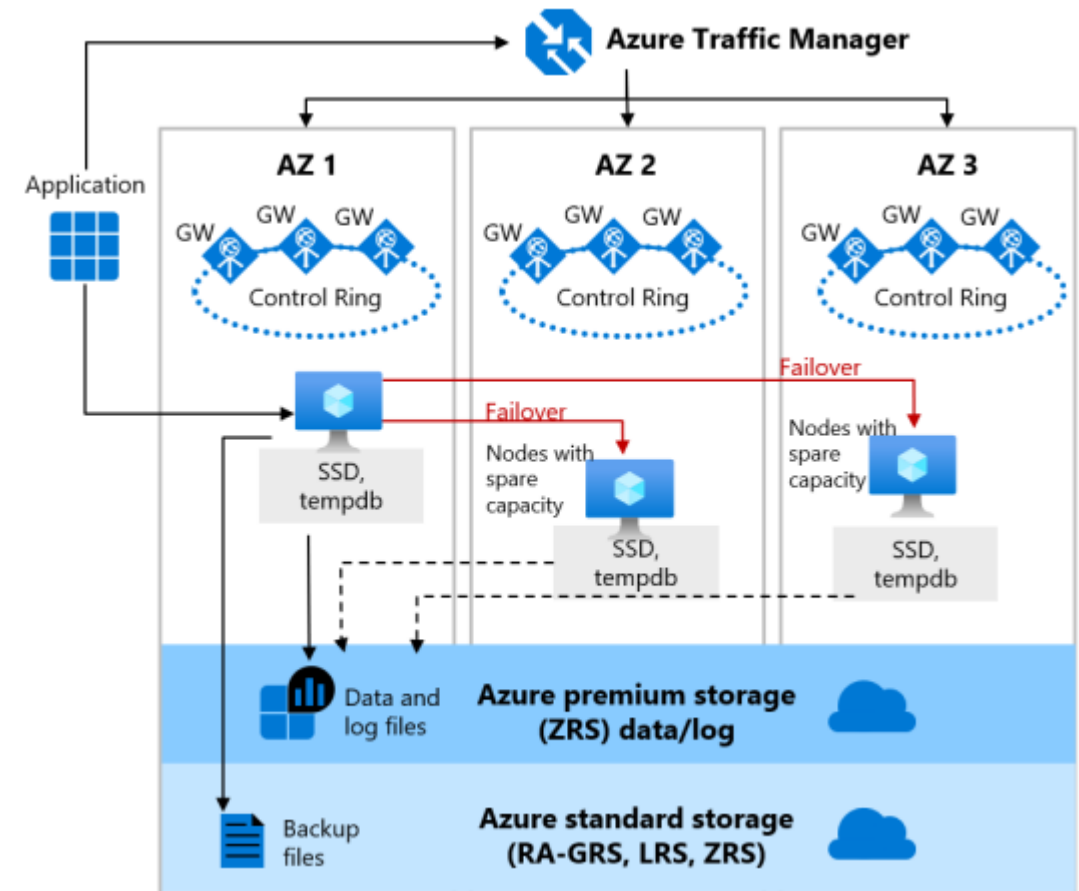
Zone redundant configuration – General Purpose

Zone redundant configurations for the General-Purpose service tier became Generally Available in December 2022. Only available in select regions.

A stateful data layer with the database files (.mdf/.ldf) that are stored in Zone Redundant Storage (ZRS) and are synchronously copied across three physically isolated Azure availability zones.

For zone-redundant serverless and provisioned General-Purpose databases, nodes with spare capacity are readily available in other Availability Zones for failover.

Zone redundancy for the General-Purpose service tier does come with an additional cost that includes charges for the extra resources required to maintain nodes with spare capacity across multiple Availability Zone




Read Scale-Out

Each database in the Premium and Business Critical service tiers is automatically provisioned with several AlwaysON replicas to support the availability SLA. These replicas are provisioned with the same compute size as the read-write.

Read Scale-Out redirects the read-only client connections to one of the ready-only replicas available instead of sharing the read-write replica.

Effectively isolate the read-only workload from the main read-write workload and doubles the compute capacity of the database or elastic pool at no additional charge.

This is ideal to scale-out for complex analytical workloads without affecting the primary OLTP workload.



How to use Read Scale-Out?

Default Configuration

- **Enabled** in **Managed Instance** Business Critical tier.
- **Disabled** in **database** placed on **SQL Database server** Premium and Business Critical tiers.

Setup Methods

- **Azure Portal**
 - Settings > Configure > Premium/Business Critical tier > Read scale-out.
- **PowerShell**
 - Set-AzSqlDatabase or
 - New-AzSqlDatabase cmdlets.
- **Azure Resource Manager REST API**
 - Create or
 - Update method

Connection


- Applications will be directed to either the read-write replica or to a read-only replica according to the **ApplicationIntent** property configured in the application's **connection string**.
- Use **ApplicationIntent=ReadOnly;** to connect to the read-only replica.

If your database is geo-replicated, be sure the read scale-out is enabled on both primary and geo-replicated secondary databases.

Backup and restore

Auto backups and Point in Time Restore (PITR)

- Full Database backup once a week
- Differential Backups every 12-24 hours
- Log Backups every 5-10 minutes
- Backup files on Azure storage with RA-GRS replicated
 - Can optionally select LRS or ZRS
- Backup Integrity checks
- Restore to new database
- Long-term retention (up to 10 years) of backups
- Geo-restore of databases if primary region down
- Restore backups of deleted databases



[Automatic, geo-redundant backups - Azure SQL Managed Instance | Microsoft Learn](#)

Setting Backup Policies

Home > Resource groups > jdSQLRG > jdsqlazure

jdsqlazure | Backups ☆ ...

SQL server

Search

Data management

- Backups
- Deleted databases
- Failover groups
- Import/Export history

Security

- Networking
- Microsoft Defender for Cloud
- Transparent Data Encryption
- Identity
- Auditing

Intelligent Performance

Available backups

Retention policies

Configure and manage your automated backup retention policies. Long-term retention policies enable you to keep full backups for up to 10 years.

Search for a database

Databases in the Basic tier are limited to a 7 day retention policy.

Database	PITR	Differential backup frequency
jdsqldb	7 Days	24 Hours

Configure policies

SQL server

Point-in-time-restore

Specify how long you want to keep your point-in-time backups. [Learn more](#)

How many days would you like PITR backups to be kept? 7

Differential backup frequency

Specify how often you want differential backups to be taken. [Learn more](#)

Take a differential backup every:

24 Hours

Long-term retention

Specify how long you want to keep your long-term retention backups. You may choose to keep yearly backups for up to 10 years. [Learn more](#)

Weekly LTR Backups

Keep weekly backups for:

6 Week(s)

Database availability and consistency

Availability

- You cannot set OFFLINE and EMERGENCY
- RESTRICTED_USER access allowed
- Dedicated Admin Connection (DAC) allowed
- Accelerated Database Recovery on by default

Consistency

- Multiple copies of data and backups
- Users can execute DBCC CHECKDB (no repair)
- Database CHECKSUM on by default
- Auto Page Repair when possible
- Data integrity error alert monitoring
- Backup and restore integrity checks
- "lost write" and "stale read" detection
- Repair without notification if no impact
- Proactive notification to customers

Accelerated Database Recovery (ADR)



- How it works
 - Uses a Persisted Version Store (PVS)
 - Independent of locking and isolation levels
 - Rollback faster than you can react
 - Undo recovery faster than you can look it up
 - Transaction log truncation not tied to active transactions



- Key questions
 - Does it require more space?
 - Will it affect performance?
 - Will I still see versions in tempdb?
 - How does it work with HA?



- Read the [paper](#)
- John's [video](#)

Demonstration

Enable and disable Read Scale-Out

- Enabling a database with read scale-out.
- Connecting to a Read Scale-Out replica.
- Disabling read scale-out.



Questions?



